

OpenStack-Health dashboard and Dealing with Data from the Gate

Matthew Treinish
mtreinish@kortar.org
mtreinish on Freenode

April 25, 2015

<https://github.com/mtreinish/openstack-health-presentation>

The OpenStack Gate

The Size of the Gate


- ▶ 12,000 jobs run in check and gate daily
- ▶

Existing Data Sources

- ▶ logstash
- ▶ subunit2sql
- ▶ graphite and grafana
- ▶ elastic-recheck
- ▶ log server
- ▶

What is OpenStack-Health

OpenStack-Health Architecture




API Server

- ▶ Wraps subunit2sql DB API using flask
- ▶
- ▶ Runs at <http://health.openstack.org>
- ▶ Continuously Deployed on every commit

subunit2sql



subunit2sql Data Collection



Frontend

- ▶ Built using AngularJS and NVD3
- ▶ Calls API server for data and renders in browser
- ▶
- ▶ Continuously Deployed on Every Commit

Current Limitations

- ▶ Only data from gate and periodic queues
- ▶ Only catches failures without subunit data (or that occur before or after subunit generation)
- ▶

Next Steps

- ▶ Include other data sources
- ▶ Use zuul for

Where to get more information

- ▶ openstack-dev ML openstack-dev@lists.openstack.org
- ▶ #openstack-qa on Freenode
- ▶ <http://git.openstack.org/cgit/openstack/openstack-health/>
- ▶ <http://git.openstack.org/cgit/openstack-infra/subunit2sql>

Questions?